

# CSE 430 Homework 3

Ryan Dougherty

## Question 5.3

What is the meaning of the term *busy waiting*? What other kinds of waiting are there in an operating system? Can busy waiting be avoided altogether? Explain your answer. When a process waits for a condition to be satisfied in a loop without giving up the processor, that is called busy waiting.

Another kind of waiting in an OS is for a processor to block on some condition, and will awake at some future time (probably by some external (to the processor) signal).

Busy waiting can be avoided altogether. However, there is an overhead in making a process sleep and/or wake.

## Question 5.5

Show that, if the `wait()` and `signal()` semaphore operations are not executed atomically, then mutual exclusion may be violated. Assume `wait()` and `signal()` are not executed atomically. The function of `wait()` is to decrement a semaphore's value atomically (in theory). Consider the case when 2 `wait()` operations execute on the same semaphore with a positive value. Since we assumed `wait()` is not executed atomically, the 2 operations may decrement the same value before an update of the value, which violates mutual exclusion. The answer for `signal()` has a similar construction.

## Question 5.9

The first known correct software solution to the critical-section problem for  $n$  processes with a lower bound on waiting  $n-1$  turns was presented by Eisenberg and McGuire... Prove that the algorithm satisfies all three requirements for the critical-section problem. We need to show that the following are true: (1) mutual exclusion is preserved, (2) the progress requirement is satisfied, and (3) the bounded-waiting requirement is met. Assume that  $P_i$  means processor  $i$ . Let CS mean "critical section".

To prove (1), see that  $P_i$  enters its CS  $\iff$  `flag[j]` is not equal to `in_cs` for all  $j$  different than  $i$ . Since only  $P_i$  can set `flag[i]` to be equal to `in_cs`, and since  $P_i$  inspects `flag[j]` only

while `flag[i]` is equal to `in_cs`, we can see that mutual exclusion is preserved.

To prove (2), see that only when a process enters its CS (and when it leaves its CS) does the value of `turn` change. Therefore, if no process executes or leaves its CS, the value of `turn` remains unchanged. The first process (that is contending to enter the CS) in the ordering of `turn, turn+1, ..., n-1, 0, ..., turn-1` will enter the CS. Therefore, the progress requirement is satisfied.

To prove (3), see that when a process leaves the CS, it must designate as its unique successor the first contending process in that same ordering as described in (2). This ordering ensures that any process wanting to enter its CS will do so within `n-1` turns. Therefore, the bounded-waiting requirement is met.

All three conditions are met, so our proof is done.

## Question 5.11

Explain why interrupts are not appropriate for implementing synchronization primitives in multiprocessor systems. Interrupts are not appropriate for implementing synchronization primitives in multiprocessor systems because disabling interrupts only prevents other processes from executing on the disabled-interrupts processor. Processes can execute anything on others. Therefore, the disable-interrupts process cannot guarantee mutual exclusion (for program state, etc.).

## Question 5.14

Describe how the `compare_and_swap()` instruction can be used to provide mutual exclusion that satisfies the bounded-waiting requirement.

```
// This code is an example of using
// compare_and_swap() to satisfy bounded-waiting.
int lock; // initially set to 0
int waiting[n]; // all initially set to 0
// ...
do {
    waiting[i] = 1;
    int key = 1;
    while (waiting[i] == 1 && key == 1)
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = 0;
    /* critical section */
    j = (i+1) % n;
    while ((j != i) && waiting[j] == 0)
```

```

        j = (j+1) % n;
    if (j == i)
        lock = 0;
    else
        waiting[j] = 0;
    /* remainder section */
} while (true);

```

Note: We can only enter the CS if `waiting[i]` or `key` is `==` to 0. The value of `key` can only be 0 if `compare_and_swap()` is executed.

Let  $P_i$  be the first process to execute `compare_and_swap()`. It will give a result of `key == 0` after the method call finishes, and therefore exits the while loop. All other processes must wait. The value of `waiting[i]` can only be 0 if another process leaves its CS. Since only `waiting[i]` is 0 for one  $i$ , the mutual-exclusion requirement is met.

For progress, we have the same argument as for mutual exclusion. A process exiting the CS either sets `lock` to 0 or `waiting[j]` to 0. Therefore, the progress requirement is met.

For bounded-waiting, when a process leaves its CS, the waiting array does a cyclic ordering from  $i+1, i+2, \dots, n-1, 0, \dots, i-1$ . It says that the first process in the ordering that is in the entry section (`waiting[j] == 1`) as the next to enter the CS. Therefore, any process wanting to enter the CS will have to wait  $n-1$  turns, and the bounded-waiting requirement is met.

## Question 5.28

Discuss the tradeoff between fairness and throughput of operations in the readers-writers problem. Propose a method for solving the readers-writers problem without causing starvation. Fairness and throughput of operations in the readers-writers problem are inversely related. Throughput can increase by favoring multiple readers (who do not modify program state) instead of favoring a single writer (for accessing and modifying shared data/values). However, doing this could result in writer starvation (i.e. can never modify data/values).

Keeping timestamps on processes that are waiting is one way to avoid starvation in the readers-writers problem. When a writer finishes writing, it wakes up the longest-waiting process. When a reader arrives and notices that other reader(s) are accessing the shared data/values, then it would enter its CS only if there are no waiting writers. Therefore, fairness is guaranteed.